

## ✓ Credit Card Fraud Detection

```

1  #Importing necessary algorithms
2
3  import pandas as pd
4  import matplotlib.pyplot as plt
5  import seaborn as sns
6
7  from sklearn.metrics import classification_report, accuracy_score
8  from sklearn.ensemble import IsolationForest
9  from sklearn.neighbors import LocalOutlierFactor

```

## ✓ Step 1 : Dataset Loading and Analysis

```

1  # Load Dataset into a dataframe using pandas
2
3  df = pd.read_csv('creditcard.csv')

1  print("Shape of the Dataset: ", df.shape) # number of rows and columns in our dataset
2  print("\n\n", df.columns) # columns/features in our Dataset

```

➦ Shape of the Dataset: (284807, 31)

```

Index(['Time', 'V1', 'V2', 'V3', 'V4', 'V5', 'V6', 'V7', 'V8', 'V9', 'V10',
      'V11', 'V12', 'V13', 'V14', 'V15', 'V16', 'V17', 'V18', 'V19', 'V20',
      'V21', 'V22', 'V23', 'V24', 'V25', 'V26', 'V27', 'V28', 'Amount',
      'Class'],
      dtype='object')

```

```
1 df.head() # first five records
```

➦

	Time	V1	V2	V3	V4	V5	V6	V7	V8
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.098698
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.085102
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	0.247676
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.377436
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.270533

5 rows × 31 columns

```
1 df.tail() # last five records
```

➦

	Time	V1	V2	V3	V4	V5	V6	V7
284802	172786.0	-11.881118	10.071785	-9.834783	-2.066656	-5.364473	-2.606837	-4.918215
284803	172787.0	-0.732789	-0.055080	2.035030	-0.738589	0.868229	1.058415	0.024330
284804	172788.0	1.919565	-0.301254	-3.249640	-0.557828	2.630515	3.031260	-0.296827
284805	172788.0	-0.240440	0.530483	0.702510	0.689799	-0.377961	0.623708	-0.686180
284806	172792.0	-0.533413	-0.189733	0.703337	-0.506271	-0.012546	-0.649617	1.577006

5 rows × 31 columns

```

1 # Print the shape of the dataframe
2
3 df = df.sample(frac = 0.3, random_state = 42) # using 30% of our dataset for next steps
4 print("Shape of the Dataset: ", df.shape)

```

➦ Shape of the Dataset: (85442, 31)

```

1 # Determine number of fraud cases in Dataset
2
3 Fraud = df[df['Class'] == 1]
4 Valid = df[df['Class'] == 0]
5
6 outlier_fraction = (len(Fraud)/float(len(Valid)))
7 print("Outlier_fraction: {0} %".format(outlier_fraction*100))
8
9 print('Fraud Cases: {}'.format(len(df[df['Class'] == 1])))
10 print('Valid Transactions: {}'.format(len(df[df['Class'] == 0])))

```

```

Outlier_fraction: 0.15942606616181745 %
Fraud Cases: 136
Valid Transactions: 85306

```

```

1 # The columns have been encrypted using PCA Dimensionality reduction to protect user identities and sensitive features
2
3 print("Description of the Dataset: ", df.describe())

```

```

Description of the Dataset:
count    85442.000000    85442.000000    85442.000000    85442.000000    85442.000000    V3    V4 \
mean     94967.874862      0.003465      0.005440     -0.004776      0.002485
std      47520.526676      1.953426      1.611981      1.520529      1.413738
min         0.000000     -37.558067     -48.060856     -33.680984     -5.600607
25%       54257.500000     -0.918861     -0.597627     -0.898758     -0.845937
50%       84962.000000      0.027558      0.063774      0.172342     -0.015543
75%      139498.000000      1.318759      0.804713      1.024434      0.748582
max      172787.000000      2.439207      21.467203      9.382558     12.699542

count    85442.000000    85442.000000    85442.000000    85442.000000    85442.000000    V5    V6 \
mean     -0.001153      0.004429     -0.006112      0.002718      0.000754
std       1.349284      1.319968      1.210313      1.208154      1.102415
min      -35.182120     -20.869626     -41.506796     -50.420090     -13.434066
25%       -0.696577     -0.768914     -0.559112     -0.210279     -0.645266
50%       -0.049751     -0.274150      0.033633      0.022630     -0.052660
75%        0.616161      0.408217      0.570121      0.328653      0.596925
max       29.016124     21.550496     36.877368     19.168327     15.594995

count    ...    V21    V22    V23    V24 \
mean    ...    0.000718    0.004176    0.000418    0.001474
std     ...    0.741520    0.726443    0.603298    0.605319
min     ...   -22.889347   -8.887017   -32.828995   -2.824849
25%     ...   -0.227053   -0.540678   -0.163221   -0.354200
50%     ...   -0.028621    0.010637   -0.012297    0.041341
75%     ...    0.187034    0.534284    0.147358    0.440782
max     ...    27.202839    8.361985    22.083545    3.990646

count    V25    V26    V27    V28    Amount \
mean     0.002232   -0.001562    0.001341    0.000050    87.300225
std      0.518783    0.481643    0.397284    0.331885    229.571240
min     -8.696627   -2.068561   -22.565679   -11.710896    0.000000
25%     -0.315966   -0.327504   -0.070385   -0.052708    5.550000
50%      0.020395   -0.055808    0.001562    0.010932    22.000000
75%      0.351311    0.237731    0.090261    0.077683    76.987500
max      6.070850    3.463246    9.200883    16.129609   10000.000000

Class
count    85442.000000
mean      0.001592
std       0.039865
min       0.000000
25%       0.000000
50%       0.000000
75%       0.000000
max       1.000000

```

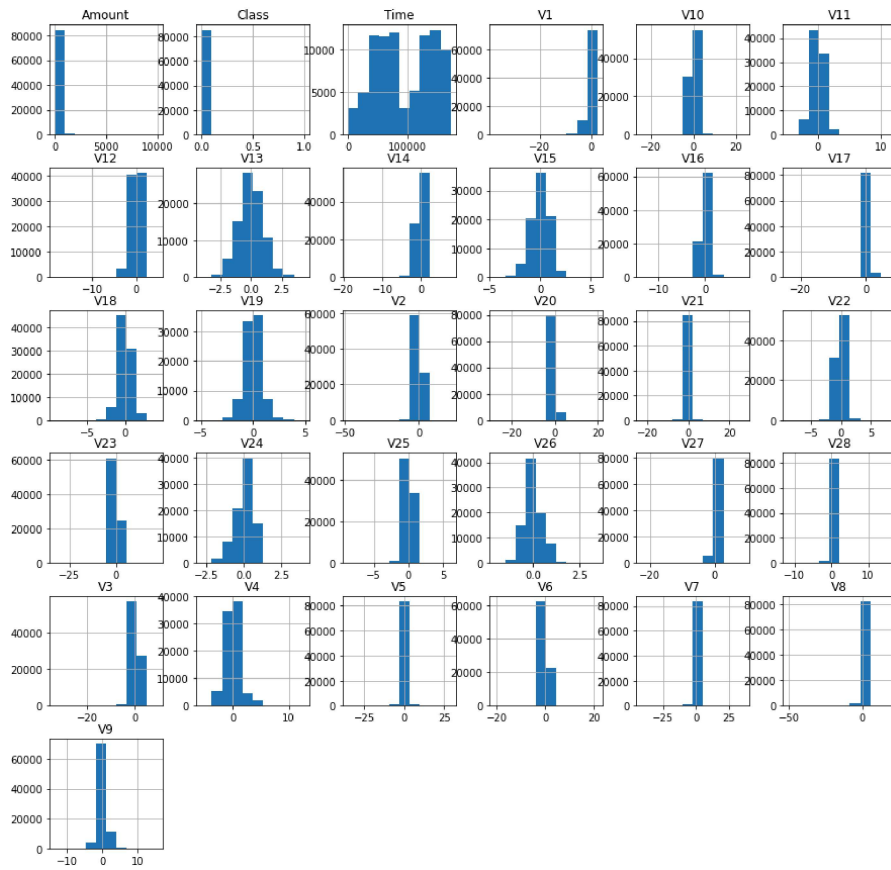
[8 rows x 31 columns]

## Step 2 : Data Visualization

```

1 # Plot histograms for each parameter
2
3 df.hist(figsize = (15, 15))
4 plt.show()


```



```

1 # Correlation matrix
2
3 corrmat = df.corr()
4 fig = plt.figure(figsize = (15, 15))
5
6 #Plotting a heatmap to visualize the correlation matrix and see features
7 # with strong correlation to the target class
8 sns.heatmap(corrmat, vmax = .6, square = True) # vmax is the max and min value you want to have for the scale
9 plt.show()

```



```
1 corrmat['Class']
```



```
Time      -0.008402
V1        -0.121864
V2         0.105051
V3        -0.208042
V4         0.128095
V5        -0.118543
V6        -0.038185
V7        -0.217359
V8        -0.001888
V9        -0.099826
V10       -0.228272
V11        0.153159
V12       -0.256577
V13       -0.012188
V14       -0.294882
V15       -0.006388
V16       -0.205082
V17       -0.345739
V18       -0.121619
V19        0.034440
V20        0.025939
V21        0.052617
V22       -0.003520
V23        0.011271
V24       -0.006823
V25        0.005641
V26       -0.002010
V27        0.013635
V28        0.007974
Amount     0.009849
Class      1.000000
Name: Class, dtype: float64
```

```
1 len(corrmat['Class'])
```

```
31
```

### Step 3 : Feature Selection

```
1 # getting columns which have correlation score > 0.01 or < -0.01, you can chose a different constant and experiment
2
3 cols = corrmat.keys()
4 cols_to_keep = []
5
6 for i in range(len(corrmat)):
7
8     if abs(corrmat['Class'][i]) > 0.01:
9
10         cols_to_keep.append(cols[i])
```

```
1 len(cols_to_keep) # the final features list we wish to keep
```

```
22
```

```
1 cols_to_keep
```

```
['V1',
 'V2',
 'V3',
 'V4',
 'V5',
 'V6',
 'V7',
 'V9',
 'V10',
 'V11',
 'V12',
 'V13',
 'V14',
 'V16',
 'V17',
 'V18',
 'V19',
 'V20',
 'V21',
 'V23',
 'V27',
 'Class']
```

```
1 # removing the 'Class' column from the features list, as it is the variable we wish to predict
2
3 cols = cols_to_keep[:-1]
```

```
1 features = df[cols] # records of all transactions, excluding the target class
2 target = df["Class"] # records of the corresponding label for each record
3
4 print(features.shape)
5 print(target.shape)
```

```
(85442, 21)
(85442,)
```

### Step 4 : Model Training and Selection

The machine learning algorithms that we are going to use are:

Local Outlier Factor Isolation Forest We are using them by importing them directly from scikit-learn; however, if you're interested in learning about their theory, you can refer to the two resources mentioned below:

- LOF: <https://towardsdatascience.com/local-outlier-factor-for-anomaly-detection-cc0c770d2ebe>

IF: [https://medium.com/@often\\_weird/isolation-forest-algorithm-for-anomaly-detection-f88af2d5518d](https://medium.com/@often_weird/isolation-forest-algorithm-for-anomaly-detection-f88af2d5518d)

```

1 # define random states
2 state = 1
3
4 # define outlier detection tools to be compared
5 classifiers = {
6     "IF": IsolationForest(max_samples = len(features),
7                           contamination = outlier_fraction,
8                           random_state = state),
9     "LOF": LocalOutlierFactor(
10         n_neighbors = 20,
11         contamination = outlier_fraction)}

1 # skipping the train, test split step because we wish the model to overfit on these features and learn
2 # a mathematical function to map the features
3
4 n_outliers = len(Fraud)
5
6 # Fit the model
7 for i, (clf_name, clf) in enumerate(classifiers.items()):
8
9     # fit the dataframe and tag outliers
10    if clf_name == "LOF":
11
12        y_pred = clf.fit_predict(features)
13        scores_pred = clf.negative_outlier_factor_
14
15    else:
16
17        # train/fit classifier on our features
18        clf.fit(features)
19        # generate predictions
20        scores_pred = clf.decision_function(features)
21        y_pred = clf.predict(features)
22
23    # Reshape the prediction values to 0 for valid, 1 for fraud.
24
25    y_pred[y_pred == 1] = 0
26    y_pred[y_pred == -1] = 1
27
28    n_errors = (y_pred != target).sum()
29
30    # Run classification metrics
31    print('Classifier {0}: \nNumber of Errors: {1}'.format(clf_name, n_errors))
32    print('Accuracy: {0}%\n'.format(accuracy_score(target, y_pred)*100))
33    print(classification_report(target, y_pred))

```



Classifier IF:

Number of Errors: 173

Accuracy: 99.797523466211%

	precision	recall	f1-score	support
0	1.00	1.00	1.00	85306
1	0.36	0.37	0.37	136
accuracy			1.00	85442
macro avg	0.68	0.68	0.68	85442
weighted avg	1.00	1.00	1.00	85442

Classifier LOF:

Number of Errors: 273

Accuracy: 99.68048500737342%

	precision	recall	f1-score	support
0	1.00	1.00	1.00	85306
1	0.00	0.00	0.00	136
accuracy			1.00	85442
macro avg	0.50	0.50	0.50	85442
weighted avg	1.00	1.00	1.00	85442